

## Problem

- ▶ Exascale nodes expected to have complex, heterogeneous memory systems
- ▶ Users won't want to learn low-level technical details
- ▶ Heterogeneity is the enemy of portability
- ▶ Middleware requires unified low-level abstraction

## Proposed Solution

Two-level library:

- ▶ Low-level interface:
  - ▶ Homogeneous abstraction over different memory devices
  - ▶ Classifies devices according to designer intent
  - ▶ Unified method of querying, allocating, and freeing
  - ▶ Intended for advanced users or middleware developers
- ▶ High-level interfaces:
  - ▶ Expose memory systems in intentional terms, rather than technical
  - ▶ Choose where to allocate, when to migrate
  - ▶ Many possible implementations
  - ▶ Basis for other research into memory management

## Project Team

Point of Contact: Mike Lang (mlang@lanl.gov)  
 Roger Pearce (pearce7@llnl.gov)  
 Maya Gokhale (gokhale2@llnl.gov)  
 Simon Hammond (sdhammo@sandia.gov)  
 Sean Williams (swilliams@newmexicoconsortium.org)  
 Douglas Otstott (douglas.otstott@gmail.com)  
 Latchesar Ionkov (lionkov@lanl.gov)  
 Ada Gavrilovska (ada@cc.gatech.edu)  
 Greg Eisenhauer (eisen@cc.gatech.edu)  
 Thaleia-Dimitra Doudali (doudali@gmail.com)  
 Terry Jones (trjones@ornl.gov)  
 Aaron Welch (welchda@ornl.gov)

## Preliminary Results

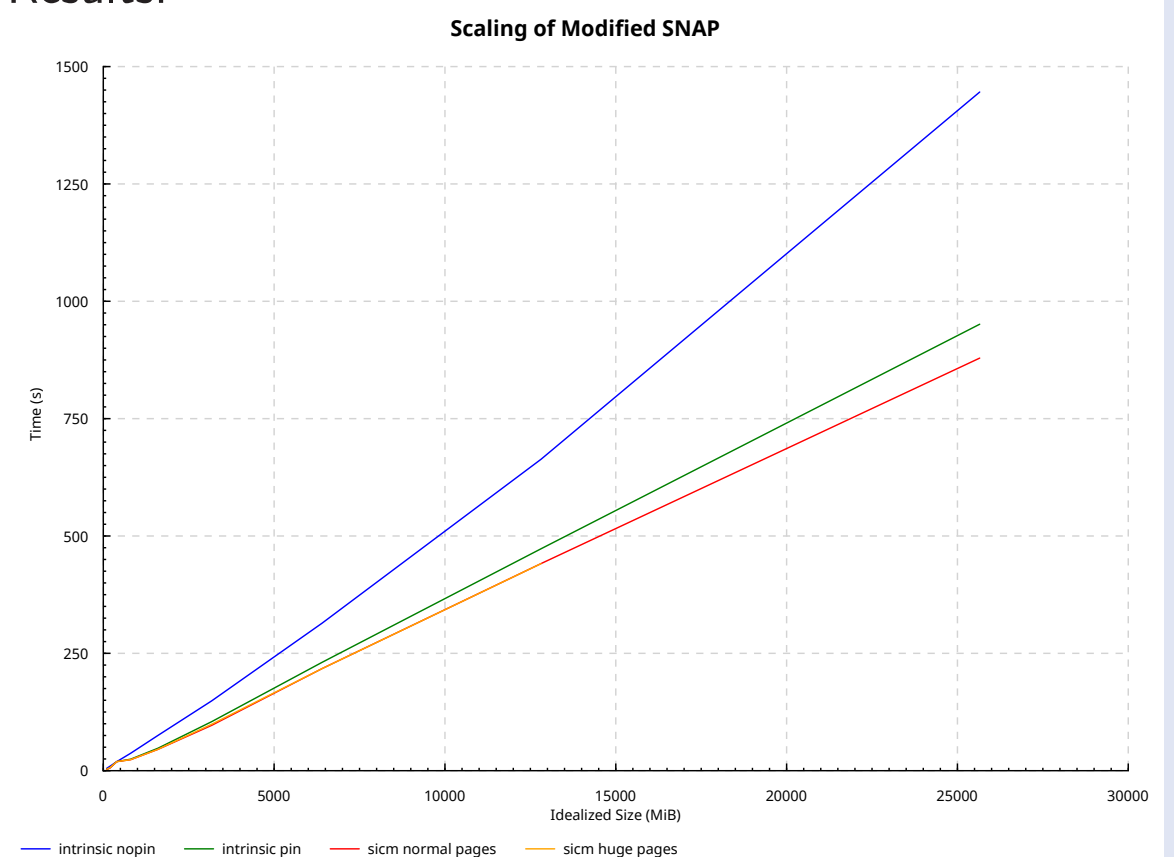
Early low-level interface:

- ▶ Currently provides abstraction over NUMA nodes, page size, and Intel Knights Landing high-bandwidth memory
- ▶ Written in C, with FORTRAN bindings

Added support to a proxy application:

- ▶ We use SNAP, a proxy particle transport simulation
- ▶ Tested NUMA pinning and huge pages

Results:



- ▶ Normal run, allocations with FORTRAN ALLOCATE
- ▶ Pin to single NUMA node, uses ALLOCATE
- ▶ Pin to single NUMA node, uses SICM
- ▶ Pin to single NUMA node, huge pages with SICM

## Repository

<https://github.com/lanl/SICM>